

CSE17040 - Distance Measures

July 28, 2020

1 MLDM Lab 2

- CB.EN.U4CSE17040

```
[1]: from scipy.spatial.distance import hamming
from scipy.spatial.distance import euclidean
from scipy.spatial.distance import cityblock
from scipy.spatial import minkowski_distance as minkowski
from math import *

import numpy as np
import pandas as pd
from matplotlib import pyplot as plt
import seaborn as sns
from random import sample
from sklearn.preprocessing import LabelEncoder
from sklearn.preprocessing import MinMaxScaler
```

1.1 Hamming Distance

```
[2]: def hamming_distance(a, b):
      return sum(abs(item1 - item2) for item1, item2 in zip(a, b)) / len(a)
```

```
[3]: row1 = [0, 0, 0, 0, 0, 1]
row2 = [0, 0, 0, 0, 1, 0]
hamming_distance(row1, row2)
```

```
[3]: 0.3333333333333333
```

```
[4]: hamming(row1, row2)
```

```
[4]: 0.3333333333333333
```

1.2 Euclidean Distance

```
[5]: def euclidean_distance(a, b):  
      return sqrt(sum((item1-item2)**2 for item1, item2 in zip(a,b)))
```

```
[6]: row1 = [10, 20, 15, 10, 5]  
      row2 = [12, 24, 18, 8, 7]  
      euclidean_distance(row1, row2)
```

```
[6]: 6.082762530298219
```

```
[7]: euclidean(row1, row2)
```

```
[7]: 6.082762530298219
```

1.3 Manhattan Distance

```
[8]: def manhattan_distance(a, b):  
      return sum(abs(item1-item2) for item1, item2 in zip(a,b))
```

```
[9]: manhattan_distance(row1, row2)
```

```
[9]: 13
```

```
[10]: cityblock(row1, row2)
```

```
[10]: 13
```

1.4 Minkowski Distance

```
[11]: def minkowski_distance(a, b, p):  
      return sum(abs(e1-e2)**p for e1, e2 in zip(a,b))**(1/p)
```

```
[12]: minkowski_distance(row1, row2, 1)
```

```
[12]: 13.0
```

```
[13]: minkowski_distance(row1, row2, 2)
```

```
[13]: 6.082762530298219
```

```
[14]: minkowski(row1, row2, 1)
```

```
[14]: 13.0
```

```
[15]: minkowski(row1, row2, 2)
```

```
[15]: 6.082762530298219
```

1.5 Cosine Similarity

```
[16]: def square_rooted(x):  
      return round(sqrt(sum([a*a for a in x])),3)
```

```
[17]: def cosine_similarity(x,y):  
      numerator = sum(a*b for a,b in zip(x,y))  
      denominator = square_rooted(x)*square_rooted(y)  
      return round(numerator/float(denominator),4)
```

```
[18]: cosine_similarity(row1, row2)
```

```
[18]: 0.9932
```

1.6 Jaccard Similarity

```
[19]: def jaccard_similarity(x,y):  
      intersection_cardinality = len(set.intersection(*[set(x), set(y)]))  
      union_cardinality = len(set.union(*[set(x), set(y)]))  
      return intersection_cardinality/float(union_cardinality)
```

```
[20]: jaccard_similarity(row1, row2)
```

```
[20]: 0.0
```

```
[21]: jaccard_similarity([0,1,2,5,6],[0,2,3,5,7,9])
```

```
[21]: 0.375
```

2 Loan Status

```
[22]: data = pd.read_csv("loan_status.csv")  
      data.head()
```

```
[22]:   grade sub_grade loan_status      purpose  
0      B         B2  Fully Paid  credit_card  
1      C         C4  Charged Off          car  
2      C         C5  Fully Paid  small_business
```

3	C	C1	Fully Paid	other
4	B	B5	Fully Paid	other

```
[23]: data.describe()
```

```
[23]:
```

	grade	sub_grade	loan_status	purpose
count	50	50	50	50
unique	6	19	2	10
top	B	B3	Fully Paid	debt_consolidation
freq	21	6	39	22

```
[24]: data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 50 entries, 0 to 49
Data columns (total 4 columns):
#   Column          Non-Null Count  Dtype
---  -
0   grade           50 non-null    object
1   sub_grade       50 non-null    object
2   loan_status     50 non-null    object
3   purpose         50 non-null    object
dtypes: object(4)
memory usage: 1.7+ KB
```

```
[25]: data.dtypes
```

```
[25]: grade           object
sub_grade          object
loan_status        object
purpose            object
dtype: object
```

```
[26]: data_crosstab = pd.crosstab(data['grade'],data['loan_status'], margins = False)
data_crosstab
```

```
[26]: loan_status  Charged Off  Fully Paid
grade
A                1           11
B                5           16
C                3            8
D                1            3
E                0            1
F                1            0
```

```
[27]: data_crosstab = pd.crosstab(data['purpose'],data['loan_status'], margins =
↪False)
```

```
data_crosstab
```

```
[27]: loan_status      Charged Off  Fully Paid
purpose
car                1             1
credit_card        0             8
debt_consolidation 4             18
home_improvement   0             1
major_purchase     1             1
medical            0             1
moving             0             1
other              4             5
small_business     1             2
wedding            0             1
```

```
[28]: data_crosstab = pd.crosstab([data['grade'],
    ↪data['purpose']],data['loan_status'], margins = False)
data_crosstab
```

```
[28]: loan_status      Charged Off  Fully Paid
grade purpose
A    credit_card        0             1
     debt_consolidation 1             7
     major_purchase     0             1
     other              0             1
     wedding            0             1
B    credit_card        0             6
     debt_consolidation 1             5
     major_purchase     1             0
     medical            0             1
     moving             0             1
     other              3             2
     small_business     0             1
C    car                1             0
     credit_card        0             1
     debt_consolidation 2             4
     home_improvement   0             1
     other              0             1
     small_business     0             1
D    debt_consolidation 0             2
     other              1             1
E    car                0             1
F    small_business     1             0
```

```
[29]: le = LabelEncoder()
df = data.copy()
```


```
df['grade'] = le.fit_transform(data['grade'])#, 'sub_grade', 'purpose'
    ↳, 'loan_status'])
df['sub_grade'] = le.fit_transform(data['sub_grade'])
df['purpose'] = le.fit_transform(data['purpose'])
df['loan_status'] = le.fit_transform(data['loan_status'])
df.head()
```

```
[29]:
```

	grade	sub_grade	loan_status	purpose
0	1	5	1	1
1	2	12	0	0
2	2	13	1	8
3	2	9	1	7
4	1	8	1	7

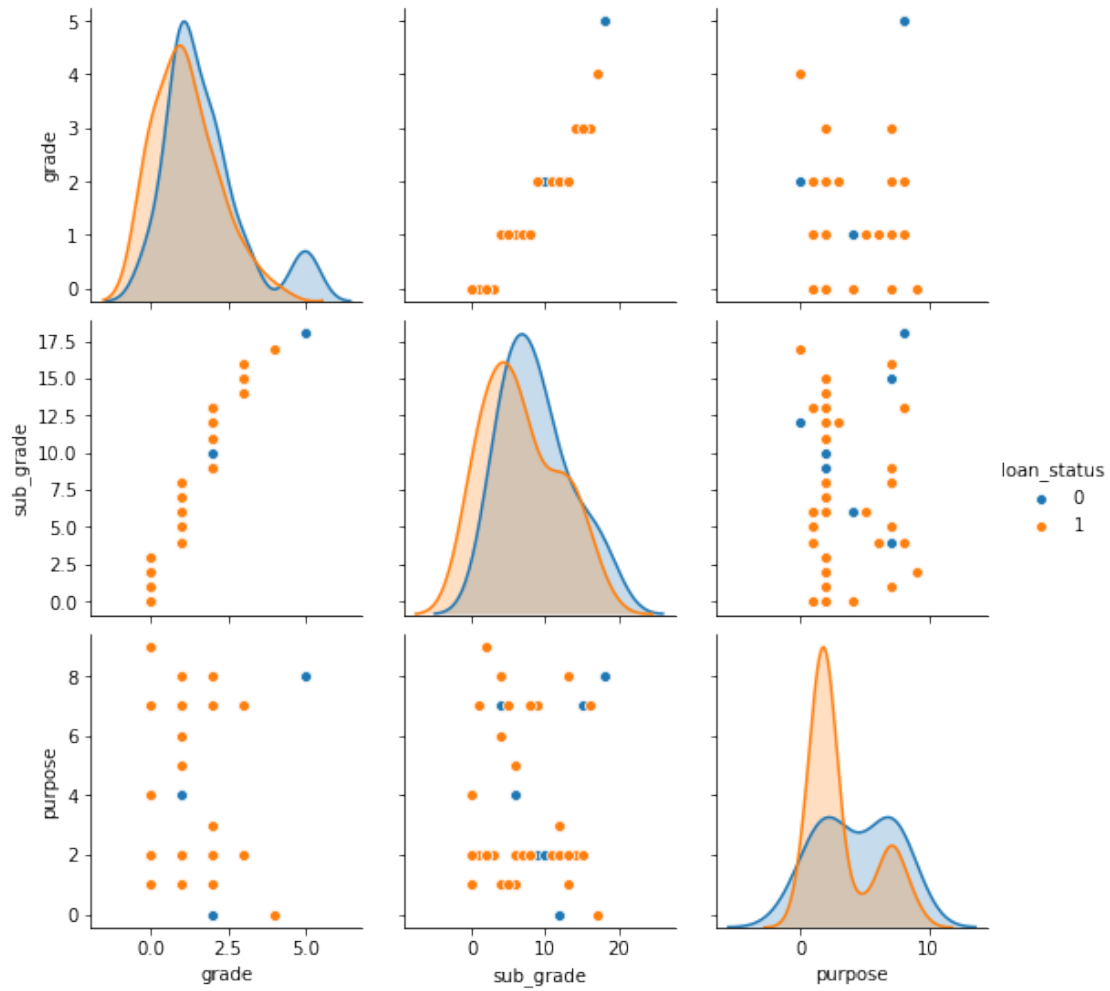
```
[30]: sns.heatmap(df.corr())
```

```
[30]: <matplotlib.axes._subplots.AxesSubplot at 0x7fcf2f18b250>
```




```
[31]: sns.pairplot(df, hue='loan_status')
```

```
[31]: <seaborn.axisgrid.PairGrid at 0x7fcf2c9cec90>
```




```
[32]: sns.boxplot(df['grade'], df['loan_status'])
```

```
[32]: <matplotlib.axes._subplots.AxesSubplot at 0x7fcf5df5ab50>
```




```
[33]: sns.boxplot(df['sub_grade'], df['loan_status'])
```

```
[33]: <matplotlib.axes._subplots.AxesSubplot at 0x7fcf2c2a1d10>
```




```
[34]: sns.boxplot(df['purpose'], df['loan_status'])
```

```
[34]: <matplotlib.axes._subplots.AxesSubplot at 0x7fcf2c076b50>
```



3 Data1.csv

```
[35]: data = pd.read_csv('Data-1.csv')
data.head()
```

```
[35]:   Country  Age  Salary  Purchased  Color
0  France  44.0   72000         No    Red
1  Spain   27.0   48000         Yes  Yellow
2  Germany 30.0   54000         No   Green
3  Spain   38.0   61000         No   Green
4  Germany 40.0     na         Yes  Yellow
```

```
[36]: data.describe()
```

```
[36]:      Age
count  114.000000
mean   39.842105
```

```
std      11.831845
min      27.000000
25%      31.000000
50%      38.000000
75%      44.000000
max      89.000000
```

```
[37]: data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 120 entries, 0 to 119
Data columns (total 5 columns):
 #   Column      Non-Null Count  Dtype
---  -
 0   Country    120 non-null   object
 1   Age        114 non-null   float64
 2   Salary     120 non-null   object
 3   Purchased  108 non-null   object
 4   Color      120 non-null   object
dtypes: float64(1), object(4)
memory usage: 4.8+ KB
```

```
[38]: data.Country.unique()
```

```
[38]: array(['France', 'Spain', 'Germany'], dtype=object)
```

```
[39]: data.isnull().sum()
```

```
[39]: Country      0
Age             6
Salary          0
Purchased      12
Color           0
dtype: int64
```

```
[40]: data['Age'].fillna(method='ffill', inplace=True)
data['Purchased'].fillna(method='ffill', inplace=True)
```

```
[41]: int(data['Salary'][0])
```

```
[41]: 72000
```

```
[42]: data.isnull().sum()
```

```
[42]: Country      0
Age             0
Salary          0
```

```
Purchased    0
Color        0
dtype: int64
```

```
[43]: for i in range(len(data['Salary'])):
      try:
          data['Salary'][i] = int(data['Salary'][i])
      except:
          data['Salary'][i] = 0
```

```
/home/vkmanojk/anaconda3/lib/python3.7/site-packages/ipykernel_launcher.py:3:
```

```
SettingWithCopyWarning:
```

```
A value is trying to be set on a copy of a slice from a DataFrame
```

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy

This is separate from the ipykernel package so we can avoid doing imports until

```
/home/vkmanojk/anaconda3/lib/python3.7/site-packages/ipykernel_launcher.py:5:
```

```
SettingWithCopyWarning:
```

```
A value is trying to be set on a copy of a slice from a DataFrame
```

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy

```
"""
```

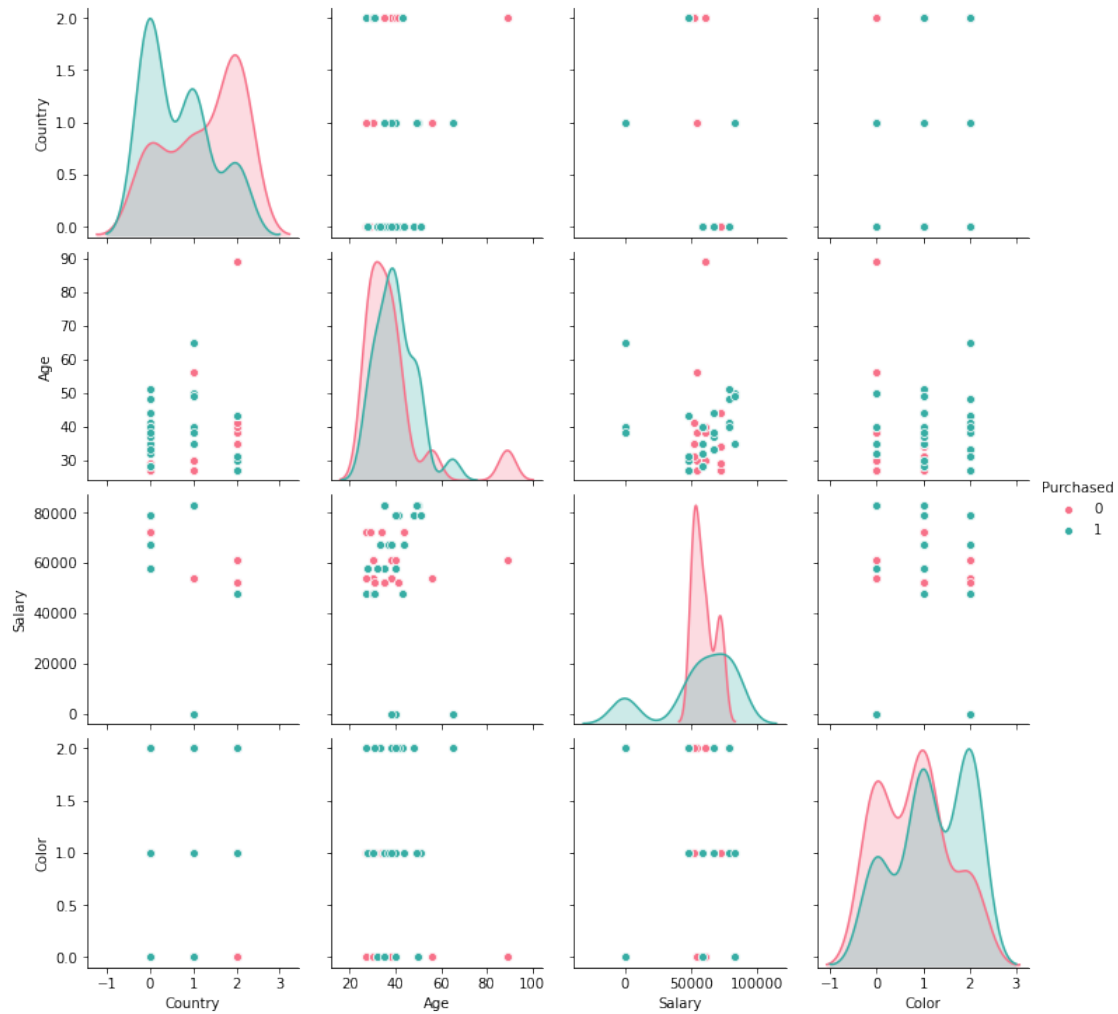
```
[44]: df = data.copy()
      df['Country'] = le.fit_transform(data['Country'])
      df['Purchased'] = le.fit_transform(data['Purchased'])
      df['Color'] = le.fit_transform(data['Color'])
      df.head()
```

```
[44]:
```

	Country	Age	Salary	Purchased	Color
0	0	44.0	72000	0	1
1	2	27.0	48000	1	2
2	1	30.0	54000	0	0
3	2	38.0	61000	0	0
4	1	40.0	0	1	2

```
[45]: sns.pairplot(df, hue = 'Purchased',palette="husl")
```

```
[45]: <seaborn.axisgrid.PairGrid at 0x7fcf2bf2ddd0>
```



3.1 Cross tab

```
[46]: pd.crosstab(data['Country'], data['Purchased'])
```

```
[46]: Purchased  No  Yes
Country
France      12  36
Germany     12  24
Spain       24  12
```

```
[47]: pd.crosstab(data['Age'], data['Purchased'])
```

```
[47]: Purchased  No  Yes
Age
```

27.0	6	3
28.0	0	3
29.0	3	0
30.0	6	3
31.0	3	3
32.0	0	3
33.0	0	3
34.0	3	0
35.0	6	6
37.0	0	3
38.0	6	6
40.0	3	12
41.0	3	3
43.0	0	3
44.0	3	3
48.0	0	3
49.0	0	3
50.0	0	6
51.0	0	3
56.0	3	0
65.0	0	3
89.0	3	0

```
[48]: pd.crosstab(data['Salary'], data['Purchased'])
```

```
[48]: Purchased  No  Yes
Salary
0           0  12
48000       0  12
52000      12   0
54000      12   0
58000       0  12
61000      12   0
67000       0  12
72000      12   0
79000       0  12
83000       0  12
```


```
[49]: pd.crosstab(data['Color'], data['Purchased'])
```

```
[49]: Purchased  No  Yes
Color
Green       18  15
Red         21  27
Yellow       9  30
```

3.2 Box plot


```
[50]: sns.boxplot(data['Country'], data['Salary'])
```

```
[50]: <matplotlib.axes._subplots.AxesSubplot at 0x7fcf2bf31550>
```




```
[51]: sns.boxplot(data['Age'], data['Salary'])
```

```
[51]: <matplotlib.axes._subplots.AxesSubplot at 0x7fcf298d8bd0>
```



```
[52]: sns.boxplot(data['Color'], data['Salary'])
```

```
[52]: <matplotlib.axes._subplots.AxesSubplot at 0x7fcf298d8550>
```



```
[53]: scaler = MinMaxScaler()
data[['Age', 'Salary']] = scaler.fit_transform(data[['Age', 'Salary']])
```

```
[54]: colors = [color for color in data['Color'].unique()]
colors
```

```
[54]: ['Red', 'Yellow', 'Green']
```

```
[55]: transformedColors = pd.get_dummies(colors)
transformedColors
```

```
[55]:
```

	Green	Red	Yellow
0	0	1	0
1	0	0	1
2	1	0	0

```
[56]: countries = [country for country in data['Country'].unique()]
countries
```

```
[56]: ['France', 'Spain', 'Germany']
```

```
[57]: transformedCountries = pd.get_dummies(countries)
transformedCountries
```

```
[57]:
```

	France	Germany	Spain
0	1	0	0
1	0	0	1
2	0	1	0

3.3 Hamming Distance

```
[58]: print('Hamming distance between Green and Red')
hamming(transformedColors.Green.values, transformedColors.Red.values)
```

Hamming distance between Green and Red

```
[58]: 0.6666666666666666
```

```
[59]: print('Hamming distance between Green and Yellow')
hamming(transformedColors.Green.values, transformedColors.Yellow.values)
```

Hamming distance between Green and Yellow

```
[59]: 0.6666666666666666
```



```
[60]: print('Hamming distance between Yellow and Red')
      hamming(transformedColors.Yellow.values,transformedColors.Red.values)
```

Hamming distance between Yellow and Red

```
[60]: 0.6666666666666666
```

```
[61]: print('Hamming distance between France and Germany')
      hamming(transformedCountries.France.values,transformedCountries.Germany.values)
```

Hamming distance between France and Germany

```
[61]: 0.6666666666666666
```

```
[62]: print('Hamming distance between France and Spain')
      hamming(transformedCountries.France.values,transformedCountries.Spain.values)
```

Hamming distance between France and Spain

```
[62]: 0.6666666666666666
```

```
[63]: print('Hamming distance between Spain and Germany')
      hamming(transformedCountries.Spain.values,transformedCountries.Germany.values)
```

Hamming distance between Spain and Germany

```
[63]: 0.6666666666666666
```

3.4 Euclidean Distance

```
[64]: purchased = data['Salary'][data['Purchased']=='Yes']
      notPurchased = data['Salary'][data['Purchased']=='No']
```

```
[65]: print("Euclidean distance between salaries of those who purchased and those who
      ↪didn't")
      euclidean_distance(purchased.values[:len(notPurchased)], notPurchased.values)
```

Euclidean distance between salaries of those who purchased and those who didn't

```
[65]: 2.2761839138049376
```

```
[66]: purchased = data['Age'][data['Purchased']=='Yes']
      notPurchased = data['Age'][data['Purchased']=='No']
```

```
[67]: print("Euclidean distance between ages of those who purchased and those who
      ↪didn't")
      euclidean_distance(purchased.values[:len(notPurchased)], notPurchased.values)
```

Euclidean distance between ages of those who purchased and those who didn't

```
[67]: 2.007270862723103
```

```
[68]: print("Euclidean distance between age and salary")
      euclidean_distance(data['Salary'].values, data['Age'].values)
```

Euclidean distance between age and salary

```
[68]: 6.421416331269408
```

3.5 Manhattan Distance

```
[69]: purchased = data['Salary'][data['Purchased']=='Yes']
      notPurchased = data['Salary'][data['Purchased']=='No']
      print("Manhattan distance between salaries of those who purchased and those who_
            ↳didn't")
      manhattan_distance(purchased.values[:len(notPurchased)], notPurchased.values)
```

Manhattan distance between salaries of those who purchased and those who didn't

```
[69]: 13.1566265060241
```

```
[70]: purchased = data['Age'][data['Purchased']=='Yes']
      notPurchased = data['Age'][data['Purchased']=='No']
      print("Manhattan distance between ages of those who purchased and those who_
            ↳didn't")
      manhattan_distance(purchased.values[:len(notPurchased)], notPurchased.values)
```

Manhattan distance between ages of those who purchased and those who didn't

```
[70]: 9.870967741935482
```

```
[71]: print("Manhattan distance between age and salary")
      manhattan_distance(data['Salary'].values, data['Age'].values)
```

Manhattan distance between age and salary

```
[71]: 67.20734551107658
```

3.6 Minkowski Distance

```
[72]: purchased = data['Salary'][data['Purchased']=='Yes']
      notPurchased = data['Salary'][data['Purchased']=='No']
      print("Minkowski distance between salaries of those who purchased and those who_
            ↳didn't with power 50")
```

```
minkowski_distance(purchased.values[:len(notPurchased)], notPurchased.values, 50)
```

Minkowski distance between salaries of those who purchased and those who didn't with power 50

```
[72]: 0.6707834228068674
```

```
[73]: purchased = data['Age'][data['Purchased']=='Yes']
notPurchased = data['Age'][data['Purchased']=='No']
print("Minkowski distance between ages of those who purchased and those who
      didn't with power 50")
minkowski_distance(purchased.values[:len(notPurchased)], notPurchased.values, 50)
```

Minkowski distance between ages of those who purchased and those who didn't with power 50

```
[73]: 0.9205480696375034
```

```
[74]: print("Minkowski distance between age and salary with power 50")
minkowski_distance(data['Salary'].values, data['Age'].values, 50)
```

Minkowski distance between age and salary with power 50

```
[74]: 0.9022134220635781
```

3.7 Cosine Similarity

```
[75]: purchased = data['Salary'][data['Purchased']=='Yes']
notPurchased = data['Salary'][data['Purchased']=='No']
print("Cosine Similarity between salaries of those who purchased and those who
      didn't")
cosine_similarity(purchased.values[:len(notPurchased)], notPurchased.values)
```

Cosine Similarity between salaries of those who purchased and those who didn't

```
[75]: 0.9016
```

```
[76]: purchased = data['Age'][data['Purchased']=='Yes']
notPurchased = data['Age'][data['Purchased']=='No']
print("Cosine Similarity between ages of those who purchased and those who
      didn't")
cosine_similarity(purchased.values[:len(notPurchased)], notPurchased.values)
```

Cosine Similarity between ages of those who purchased and those who didn't

[76]: 0.4806

```
[77]: print("Cosine similarity between age and salary")
      cosine_similarity(data['Salary'].values, data['Age'].values)
```

Cosine similarity between age and salary

[77]: 0.6874

3.8 Jaccard Similarity

```
[78]: purchased = data['Salary'][data['Purchased']=='Yes']
      notPurchased = data['Salary'][data['Purchased']=='No']
      print("Jaccard Similarity between salaries of those who purchased and those who_
            ↪didn't")
      jaccard_similarity(purchased.values[:len(notPurchased)], notPurchased.values)
```

Jaccard Similarity between salaries of those who purchased and those who didn't

[78]: 0.0

```
[79]: purchased = data['Age'][data['Purchased']=='Yes']
      notPurchased = data['Age'][data['Purchased']=='No']
      print("Jaccard Similarity between ages of those who purchased and those who_
            ↪didn't")
      jaccard_similarity(purchased.values[:len(notPurchased)], notPurchased.values)
```

Jaccard Similarity between ages of those who purchased and those who didn't

[79]: 0.36363636363636365

```
[80]: print("Jaccard similarity between age and salary")
      jaccard_similarity(data['Salary'].values, data['Age'].values)
```

Jaccard similarity between age and salary

[80]: 0.06666666666666667

[]: